

Predicting Rating Gains in maimai DX

A Player-Specific Statistical Model in the Rasmai Recommender

Cardin Nguyen

cardin.nguyen.ink

September 30, 2026

Abstract

Rasmai turns a maimai DX player's scores into a ranked list of charts to replay, ordered by the rating each replay would add. This paper documents every calculation behind that list: the official rating formula, the two-pool best-50 total, a per-player skill curve, a normal model of run-to-run variation, target selection by expected gain, self-calibration against recorded plays, note-level loss accounting, technique measures read from chart notation (*simai*), trait detection and route planning. The methods are then applied to an example, one Rasmai user's export of 554 scored charts, whose numbers are illustrative and will differ for other players. It is run through the project's own code with the chart databases and notation fetched on the same day. The rerun reproduces all 554 stored chart ratings and a best-50 total of 13,579, one below the 13,580 the site shows. After calibrating on 92 recorded plays, the model gives one replay of GIGANTØMAKHIA a 36% chance of reaching SS, which would add 4 rating points.

1. Introduction

Rasmai is a Discord bot and web dashboard (rasmai.lol) for [maimai DX](#), SEGA's arcade rhythm game. It is named after Ras, a character from the game. It reads a player's scores from maimai DX NET ([international site](#)), the game's official score site, and answers one question: which chart should this player replay tonight to raise their rating the most? The source code is at github.com/Ne-k/Rasmai.

The question has three parts. First, how does the game turn a score into rating, and which scores count? Second, what score can this particular player realistically reach on a given chart? Third, how likely is a single run to land it? The game answers only the first. The rest is Rasmai's own model, built from the player's history and checked against their later plays.

This paper explains each step in the order the code runs it. Section 2 describes the system and Section 3 gives the formulas, constants and thresholds. Section 4 applies them to a real player export recorded on 30 September 2026, and Section 5 discusses design choices and limits. Source files are linked where they are discussed, and Appendix A lists them all.

A note on terms. *Achievement* is the percentage score of a play, from 0 to 101%; rating counts it only up to 100.5%. A chart's *constant* is its hidden internal difficulty, such as 13.0 or 13.6, which the game shows only as a level like "13" or "13+". *Rating* is the integer a play earns from constant and achievement, and a player's *rating* is the sum of their best 50 chart ratings. *Simai* is the plain-text notation maimai charts are written in.

2. System Overview

Rasmai is a Python program with a separate web front end. The Python side has four layers, and the calculations in this paper live in the third.

1. **Scraping** ([scraping](#)). A signed-in session reads the player's profile, score tables, recent plays, per-play judgement pages and play counts from maimai DX NET. Other sources fill in what the score site leaves out:

- chart constants, versions and regions from [otoge-db](#);
 - revised constants, per-chart regions and song aliases from the `dxdata.json` file of [dxrating](#);
 - search aliases from [maimai-chart-browser](#);
 - English titles, videos, unlock notes and area names from [SilentBlue RemyWiki](#);
 - pattern tags written by editors, and chart notation, from [maiノーツ](#);
 - chart notation in bulk from [Maichart-Converts](#).
2. **Storage** ([storage/](#)). Accounts, scores, play history, judgement pages and rating readings are kept per user. Every play the bot has seen is stored, so history outlives the game's own 50-play list.
 3. **Engine** ([engine/](#)). Pure functions over the stored data: rating and best-50 pools, the play profile, recommendations, route planning, notation measures, and the insights (calibration, traits, judgements, forecast).
 4. **Interfaces** ([bot/](#), [web/](#)). A Discord bot renders results as embeds and posters. A Next.js dashboard calls an internal JSON API that the bot serves on 127.0.0.1:8765.

To stay light on the game's servers, a command first checks the profile and recent plays, which takes a few seconds. It rebuilds from the stored copy unless a new play, play count or rating appears. A full read happens when something changed, when the stored copy is a week old, or on `/refresh`.

The engine runs the same sequence on every request ([scraper/analyzer.py](#)):

1. **Enrich**. Match each scraped score to its chart, attach the constant, flag whether the chart belongs to the current game version, and compute its rating (`enrich_songs`).
2. **Pool**. Split ratings into the new and old pools and total the best 50 (`build_best50`).
3. **Profile**. Fit the player's skill curve, spreads and tier offsets from their scores and plays (`build_play_profile`). Inside this step the model also calibrates itself against recorded plays and fits the player's traits.
4. **Recommend**. For each chart, choose a target rank, price it, and rank charts by expected rating gain (`generate_recommendations`).

The steps depend on each other in this order, so an error in the enrichment step, such as a wrong constant, carries into every later number.

3. Methodology

3.1 Ranks and rating coefficients

Achievement maps to a named rank, and each rank carries a coefficient. The table below is the game's DX table as encoded in [rating.py](#) (`RANK_TABLE`).

Rank	Minimum achievement (%)	Coefficient
SSS+	100.5	22.4
SSS	100.0	21.6
SS+	99.5	21.1
SS	99.0	20.8
S+	98.0	20.3
S	97.0	20.0
AAA	94.0	16.8
AA	90.0	15.2

Rank	Minimum achievement (%)	Coefficient
A	80.0	13.6
BBB	75.0	12.8
BB	70.0	12.4
B	60.0	12.0
C	50.0	10.4
D	0	9.5

A score exactly 0.0001 under certain boundaries uses its own coefficient: 100.4999 uses 22.2 instead of 21.6, 99.9999 uses 21.4 instead of 21.1, 99.4999 uses 21.0, 98.9999 uses 20.6, 96.9999 uses 17.6 and 59.9999 uses 11.2 (`BORDERLINE_COEFFICIENTS`).

3.2 Rating of one chart

With constant c and achievement a , capped at 100.5, the chart's rating is the whole number below the product:

$$R(c, a) = \left\lfloor c \cdot \frac{\min(a, 100.5)}{100} \cdot k(a) \right\rfloor$$

Here $k(a)$ is the coefficient from the table above. A score that fails to parse, or any non-positive constant or achievement, is worth 0 rather than raising an error. Because the coefficient jumps at each rank boundary, rating is a staircase in achievement, and crossing a boundary is worth a full step. The reverse function `accuracy_for_rating` finds the lowest achievement that reaches a wanted rating, searching bands from the bottom and nudging by 0.0001 to defeat floor rounding.

3.3 Chart constants and versions

The score site prints only a level such as "13+". Rasmai looks the constant up in the chart database by the key (title, chart type, difficulty), folding case and falling back to a title stripped of symbols and spacing (`charts.py`). Where two songs share a title, the level printed on the site breaks the tie. A level from 7 up covers two bands: "13" is 13.0 to 13.5 and "13+" is 13.6 to 13.9. If a chart is missing from the database, its constant is estimated as the middle of its band, so a missing "13+" is read as 13.75 (`level_range`, `enrich_songs`). A raw database level of "13+" with no constant is parsed as 13.6 (`parse_constant`).

A chart is *new* when its version equals the player's current game version. That version is the highest one with at least 8 of the player's scored charts (`detect_current_version`).

3.4 The best-50 total

The rating shown on a profile is the sum of two pools (`pools.py`):

- **New pool:** the top 15 chart ratings among current-version charts.
- **Old pool:** the top 35 chart ratings among all other charts.

Utage charts are excluded. Each chart counts once, at its highest rating. In symbols, with N the new pool and O the old pool:

$$\text{Rating} = \sum_{i=1}^{15} N_{(i)} + \sum_{i=1}^{35} O_{(i)}$$

where $N(i)$ and $O(i)$ are the chart ratings sorted from highest to lowest. Each pool keeps its *cutoff* (the rating in its last slot) and its *next value* (the rating just outside it). These two numbers decide what a future play is worth.

Marginal gain. For a chart with rating r' after a new score, a pool changes as follows (`gain_if`):

$$\Delta = \begin{cases} \max(r', \text{next}) - r_{\text{now}} & \text{chart already in the pool} \\ r' - \text{cutoff} & \text{chart outside the pool and } r' > \text{cutoff} \\ 0 & \text{otherwise} \end{cases}$$

A chart inside the pool gains its whole improvement. A chart outside it gains only what it adds above the cutoff, because it displaces the chart in the last slot. The gain is measured against today's pools. Once several planned improvements land, the cutoffs rise, so the recommender also simulates the pools after all of them (`summarise`), giving a smaller, honest "reachable gain" beside the naive sum.

3.5 Rating colours

The total maps to a colour band, following [SilentBlue RemyWiki's rating page](#): white from 0, blue from 1,000, green 2,000, yellow 4,000, red 7,000, purple 10,000, bronze 12,000, silver 13,000, gold 14,000, platinum 14,500, rainbow 15,000 and kiwami 16,000. Gold and above are cut into 250-point steps, up to four per band (`rating_band`, `rating_step`).

3.6 The skill curve

The official formulas say nothing about what a player will score. Rasmai estimates that from the player's own scores, as a curve $E(c)$ of the achievement they reach at constant c when a run goes well. The fit (`profile/fit.py`, `profile/model.py`) has five steps.

1. **Bucket.** Each scored chart is placed in a bucket of half a constant, `round(2c)/2`.
2. **Summarise each bucket.** The bucket's representative score is its upper quartile when it holds at least 6 charts, and its 65th percentile when it holds fewer. This reads what the player does when they try and ignores careless runs; a thin bucket gets the more cautious percentile.
3. **Fit a line.** A [weighted least-squares](#) line $L(c)$ through the representatives, each weighted by $\min(n, 10)$, where n is the bucket's chart count. The slope is clamped to $[-2.5, -0.2]$. With fewer than 8 scores or 3 buckets, the fit falls back to $101 - 0.3c$.
4. **Measure between-chart spread.** For buckets with at least 3 charts and within 4 constants of the hardest bucket, take the upper half of the scores and subtract the bucket's representative. σ_c is the population standard deviation of those residuals, clamped to $[0.4, 2.0]$ (0.9 when there are 4 or fewer). Easy charts are left out because they are played carelessly and would inflate it.
5. **Correct locally.** A straight line misses kinks, so $L(c)$ is blended with a Gaussian [kernel average](#) of the bucket representatives r_j :

$$\hat{L}(c) = \frac{\sum_j w_j r_j}{\sum_j w_j}, \quad w_j = \min(n_j, 10) e^{-\left(\frac{c_j - c}{0.6}\right)^2}, \quad \beta = \frac{W}{W + 3}, \quad W = \sum_j w_j$$

$$E_{\text{raw}}(c) = \beta \hat{L}(c) + (1 - \beta) L(c)$$

Beta approaches 1 where the player has many charts nearby and 0 where they have none, so the curve trusts the data where there is data. Three guards then apply:

- **Dense ceiling.** Above the hardest bucket with at least 3 charts, the curve loses a further 0.8 points per constant, so a thin top end never promises high ranks.

- **Own ceiling.** Between populated buckets (4 or more charts), the curve may not exceed their interpolated value by more than 0.3.
- **Monotone smoothing.** From the bucket with the most charts, the curve may not rise for harder charts or fall for easier ones. This removes kinks that once made the model promise more at constant 15 than at 14.6.

The curve is capped at 100.5. Two summary constants come from it: the *reach* constant is the hardest c with $E(c) \geq 97$ (an S is expected), and the *comfort* constant is the hardest with $E(c) \geq 99$, never less than one level below reach.

Tier offsets. A Master 13 and an Expert 13 are not the same chart to the same hands. For each difficulty tier with at least 3 scores, the offset is the median of that tier's residuals from the curve, shrunk by $n / (n + 12)$ so a thin tier stays near zero. Offsets are centred on the player's own chart mix and clamped to ± 6 ([profile/build.py](#)). The median resists a few abandoned runs. The tier-adjusted curve used everywhere below is

$$m(c, \text{tier}) = E(c) + \text{offset}_{\text{tier}} + \text{shift} + \text{shift}_{\text{tier}}$$

where the two shifts come from calibration (Section 3.11) and are zero until there are enough recorded plays.

3.7 Two kinds of spread

The model keeps two widths because they answer different questions.

Between-chart spread (`sigma_at`) says how uncertain the score is on a chart the player has never tried, where the chart's own difficulty *for them* is unknown. It is $\max(0.3, \sigma_c)$, multiplied by $1 + 0.5(c - \text{dense ceiling})$ above the dense ceiling, raised to at least 0.8 when fewer than 40 charts are scored, multiplied by the calibration scale s of Section 3.11, and capped at 3.0.

Run-to-run spread (`run_sigma_at`) says how far one run on a chart the player already holds falls from the next. The difference between two plays of the same chart cancels the chart's difficulty, leaving only run variation. From the stored plays, Rasmai takes successive differences Δ on each chart, dropping runs more than 5 points under that chart's best as abandoned, and estimates

$$\hat{\sigma} = \frac{1.4826 \cdot \text{median} |\Delta|}{\sqrt{2}}, \quad \sigma_{\text{run}} = \text{clamp}\left(0.5 + (\hat{\sigma} - 0.5) \frac{m}{m + 15}, 0.3, 1.5\right)$$

where m is the number of pairs. The factor 1.4826 turns a [median absolute deviation](#) into a standard deviation, and dividing by $\sqrt{2}$ converts a difference of two runs into one run. Fewer than 12 pairs leaves the default of 0.5 ([insights/calibration.py](#), `run_spread`). In use, the run spread is at least 0.25, rises to the between-chart spread above the dense ceiling, and is multiplied by the calibration scale s . Using the between-chart number for a chart already scored made every rank look reachable, which is why the two are kept separate.

3.8 Expectation on a specific chart

For a chart the player has scored, `chart_expectation` starts from the tier-adjusted curve m and pulls it toward the player's own record. Let b be the best score and n the play count: the larger of the official count and the appearances in the recent-play list, and at least 1 if any score exists. Then, in order:

$$\sigma = \sigma_{\text{run}} \sqrt{1 + \frac{1}{n}}, \quad w = \frac{n}{n + 3}, \quad \mu = m - w \max(0, m - b), \quad \mu \leftarrow \max(\mu, b - 0.5\sigma), \quad \mu \leftarrow \min(\mu, b + h\sigma)$$

The square-root factor widens the spread for a chart played once, whose standing is estimated from one sample. The next term says a chart played often is expected near its best even when the curve says otherwise. The floor says a best well above the curve marks a chart the player is good at. The cap, with headroom $h = 0$ (`BEST_HEADROOM`), stops the centre from sitting above a proven score. The source reports that, over 246 recorded runs, moving h from 0.3 to 0 called the exact rank right 64% of the time instead of 61%.

One adjustment comes last. If the chart appears two or more times in the recent-play list and the latest score beats the oldest, μ rises by $\min(0.4, 0.25 \times \text{that improvement})$. Because it comes after the cap, it can lift μ up to 0.4 above the best.

Dropped bests. A best far under the curve on a chart played at most twice is probably one abandoned run, not the chart's difficulty. It is dropped when $b < m - \max(2, 2\sigma c)$, with σc the between-chart spread, and both the play count and the chart's runs in the recent-play list are at most 2. Such a chart is treated as untried: its *first-pass* estimate is $m - 0.35 - 0.4\sigma c$, with spread σc . When the chart look-up shows odds for a chart the player has never played ([planning/predict.py](#)), it widens that spread by a further factor of 1.5, because held-out first passes scattered 1.8 times as widely as repeat plays.

3.9 Odds and target ranks

The chance that one run reaches a target achievement t is the upper tail of a [normal distribution](#), where Φ is its cumulative distribution function:

$$P(a \geq t) = \Phi\left(\frac{\mu - t}{\sigma}\right)$$

For each played chart, [picks/targets.py](#) chooses one target from the ladder SSS+ (100.5), SSS (100.0), SS+ (99.5), SS (99.0), S+ (98.0), S (97.0), AAA (94.0) and AA (90.0). A rank is a candidate only if it passes every test below.

1. It is above the current score and at or below a *ceiling* of $\min(100.5, \mu + s \cdot \max(\sigma, \sigma c))$, where s is the challenge level's sigma count. Reach is judged with the wider spread because how far a chart can be taken is a question about practice, while the odds of one run use σ alone.
2. It sits within a few points of the current score a : at most $a + \min(6, 1.5 + r + s)$, where $r = \max(0, \mu - a)(1 - n/(n+3))$ is the recovery still expected on an unfamiliar chart. A dropped best is instead held to $\mu + 0.5\sigma$.
3. For S or higher, it needs *proof*: it may not exceed, by more than the level's slack, the best score the player holds in any half-level bucket at $c - 0.7$ or above (`best_near`). A player who has never reached 97% near a 12.5 is not told to S a 13.
4. Its odds meet the level's minimum, and its rating exceeds the chart's current rating.

Among the survivors, the level's rule picks one:

Level	Sigmas s	Slack	Minimum odds	Rule	Value of a rank
Easier	0.35	0.5	0.30	likely	odds + gain / 10,000
Balanced	1.0	1.0	0.16	value	gain $\times$ odds
Challenging	1.3	2.0	0.08	gain	gain + odds
Long shots	2.2	3.0	0.03	gain	gain + odds

Balanced is the default. It picks the rank that maximises the odds times the rating the chart itself would gain, not simply the next rank up.

3.10 Ranking the picks

A played chart's pick score multiplies four quantities ([picks/generate.py](#)):

$$\text{score} = \Delta R \times P \times \alpha \times \phi$$

- **ΔR** is the best-50 gain from Section 3.4, computed on the target's rating.
- **P** is the odds from Section 3.9.
- **α (affinity)** rewards charts like the ones the player picks. Recent plays give counts by chart type, genre and difficulty. Each count becomes a multiplier $1 + \gamma \tanh(\text{count} / \text{average} - 1)$, with $\gamma = 0.35$ for type, 0.25 for genre and 0.20 for difficulty, so no factor strays more than γ from 1. The three multiply together.
- **ϕ (freshness)** discounts charts already being ground, since the easy points are gone. It is 0.45 for a chart with 3 or more runs in the recent-play list; 0.75 (one run) or 0.6 (two) if its latest run is among the last 5 plays; 0.95 if among the last 15; otherwise 1.

The score orders the near misses below. The rating moves themselves are sorted by ΔR first and by odds second, because ΔR is a whole number that many charts share. Only the Easier level sorts by $\Delta R \times \text{odds}$ instead. An optional beta reading of the player's taste can multiply the odds in that tie-break; without it, the weight is exactly 1. At most 6 unplayed charts appear among the picks.

Near misses. A chart whose target would not move the pools can still be one good run from entering them. Rasmai asks what achievement gives the pool's cutoff + 1 (`accuracy_for_rating`). If that is within $2 \cdot \max(0.5, \sigma c)$ of the expected score, the chart is listed as a near miss, ranked below every real gain by $\alpha\phi / (1 + \text{gap})$.

Unplayed charts (`picks/candidates.py`). The search covers Expert, Master and Re:Master charts that are playable in the player's region and not locked, with constants from reach - 1.5 up to the *search ceiling* + 0.3. The search ceiling is the larger of the hardest constant with an S and the reach constant, so one failed attempt at a 15 does not open up 15s. The expected first pass is $m - 0.35 - 0.4\sigma c$. A chart is offered if that estimate is at least 94 (80 while its pool has open slots) and would enter the pool. Its score is $\Delta R \cdot P \cdot \alpha \cdot \text{proximity} \cdot 0.7$, where $\text{proximity} = 1 / (1 + |c - \text{comfort}| / 3)$, and a new-pool chart scores 1.4 times more while the new pool has open slots. When fewer than 12 rating moves remain, the list is continued with charts worth a first run.

3.11 Calibration against recorded plays

Every stored play made with an earlier best on record is predicted from that earlier best and then compared with what happened (`insights/calibration.py`). For each such play, $z = (a - \mu) / \sigma$, with σ at least 0.3; plays with $z < -4$ are dropped as abandoned. With at least 15 usable plays (n), the model corrects itself:

$$s = 1 + (\text{clamp}(\rho, 0.6, 1.8) - 1) \frac{n}{n + 30}, \quad \rho = 1.4826 \text{ median } |z - \text{median}(z)|$$

$$\text{shift} = \text{clamp}\left(Q_{0.65}(a - \mu) \frac{n}{n + 30}, -2, 2\right)$$

The scale s multiplies both spreads. The shift moves the curve, because the model's centre is meant to be "a good run", so the upper third of residuals should sit on zero. Each tier with at least 8 plays (m of them) also gets its own shift of $(Q_{0.65} \text{ of the tier} - Q_{0.65} \text{ of all}) \cdot m / (m + 20)$, clamped to ± 2 . Every correction is shrunk toward no change while evidence is thin. The result also reports *coverage*, the share of plays with $|z| \leq 1$ (about 68% for a well-sized normal model), and the model's predicted new-best rate against the actual rate.

3.12 Where the points went: note-level losses

A play's judgement page lists, for each note type (tap, hold, slide, touch, break), how many were critical perfect, perfect, great, good or miss. `losses.py` splits the achievement lost across those types. Each note type is worth a share of the 100% base in proportion to its weight:

Note type	Weight
Tap	1
Touch	1
Hold	2
Slide	3
Break	5

With counts n_k and weights w_k , one weight unit is worth $b = 100 / \sum w_k n_k$. A great keeps 80% of a note, a good 50%, and a miss nothing, so the loss on a non-break type is

$$\text{loss}_k = w_k b \left(\frac{\text{great}}{5} + \frac{\text{good}}{2} + \text{miss} \right)$$

Breaks are handled differently. From the base, a break good is charged 3 units and a break miss 5 units; a break great is not charged here. Each break also carries an equal share of a 1% bonus. Only a critical perfect keeps its whole slice: the code assumes a plain perfect keeps half, a great 40%, a good 30% and a miss none. The judgement page does not show finer grades inside a judgement, so these are assumptions. Whatever the counts leave unexplained, up to 101 minus the achievement, is assigned to breaks. The losses therefore add up to 101 minus the achievement by construction, and the size of that remainder shows how much the counts could not price. The bonus can be reported on its own, so that a player not hunting criticals does not look weak at breaks. The arithmetic follows [maimai-score-details](#) by SpiritsUnite (Apache-2.0).

Across many plays, [judgements.py](#) compares each type's share of the loss with its share of the points at stake. Their ratio is the *tilt*: 1.0 means a type costs what it was worth. With at least 3 plays stored, a type is called weak when its tilt is at least 1.5 and it carries at least 15% of all loss. The same pass reports the share of timing errors that were late, and the loss per play.

3.13 Techniques read from chart notation (*simai*)

Editors at [maiノーツ](#) have tagged about one chart in ten, almost all Master and above. To reach the rest, Rasmai reads each chart's *simai* notation and measures what the notes ask of the hands ([engine/simai/](#)).

Parsing ([parse.py](#)). A *simai* chart is a run of beat slots separated by commas. `(120)` sets the tempo and `{8}` sets how much of a bar each comma advances. Buttons 1 to 8 run clockwise round the ring, and touch areas are A to E. A slash joins notes struck at the same moment. The parser turns this into a list of notes, each with a time in seconds, a kind (tap, hold, slide, touch, touch hold), a position, and, for holds and slides, a duration, shape, end button and waiting time. A reading that leaves any character unexplained is refused rather than guessed at. When a file holds several charts, the one whose note count is nearest the published count is taken.

Measures ([features.py](#), [techniques.py](#)). Every measure is a rate or a share of the chart's notes, so long and short charts compare. The general ones are notes per second, the most notes in any one second (*bursts*), the share of slides that travel in 0.35 s or less, the share of touch notes landing together, the mean ring distance between consecutive notes (*reaching across the screen*), the share of notes struck while an earlier hold or slide is still running (*both hands at once*), and the count of tempo changes. The technique detectors work on struck notes in time order, since a slide's body is followed rather than struck:

- A *run* is a stretch of notes at most 0.18 s apart whose gaps stay within 35% of each other.
- A run of 3 or more on one button is a *jack*. A run of 4 or more alternating between buttons is a *trill*: *on the spot* between two buttons less than 3 apart, *across the screen* between two buttons 3 or more apart. A run

of 5 or more with every other note on one button is a *trill against a held button*. Any other run of 8 or more is a *stream*.

- A *bouncing rhythm* is a long-short-long pattern of gaps, each gap at most 0.5 s, whose long-to-short ratio is 1.6 to 3.4.
- *Spinning round the ring* is 5 or more notes, each at most 0.5 s apart, stepping at most 2 buttons the same way round.
- Slide measures count *chained slides* (a star starting where the last ended, within 1.5 s), *delayed slides* (notes struck while a star waits), *fan slides*, *a trill over a slide*, *overlapping loop slides*, *slides at different speeds* (overlapping slides whose durations differ by 2 times or more), *slides fired from one spot* (3 or more from one button within a phrase), and *a slide traced straight back*. Touch measures count *touch sweeps* (3 or more pads in a run) and *touch clusters* (3 or more pads at once).

Naming. A chart earns a measure's label when it sits in the top quarter of the game for that measure, taken over charts of level 12 and above, once at least 40 charts have been read. Three rare measures use a fixed bar instead, because three quarters of charts score zero on them: 1 tempo change, trills across the screen at 0.004, and slides at different speeds at 0.001. The labels join the chart's other attributes as tags.

Validation. A measure is named only if it passes two gates recorded in the source: it must rise with chart level, and it must run higher on charts the editors tagged for the same thing. The source reports, for example, streams with a level correlation of +0.61 and 4.1 times more common on charts tagged 乱打, and trills at +0.55 and 5.1 times on トリル. Jacks is the weakest named measure, at +0.03 and 1.9 times. Measures that fell with level, such as the old slide-shape spin count and the share of simultaneous notes, are still stored but no longer named. A scraped title is matched to its notation exactly first, then by the loose title match of Section 3.3; the source reports that the loose match reaches 95% of the charts a player has played, against 62% for the exact one ([scraping/simai.py](#)).

3.14 Traits

A *trait* is a chart attribute on which the player scores above or below their own curve ([insights/traits.py](#)). The attributes come from [insights/tags.py](#): chart type, the simai labels of Section 3.13, tempo (130 BPM or under, 210 or over), note count (under 620, 890 or more), era, genre, designer, and editors' pattern tags where they exist. Editor tags that describe fame rather than technique, such as "notorious", are left out.

1. **Observations.** Each chart's best is one observation, with residual $r = a - m$. Every recorded play is another, and the plays of one chart share one unit of weight, so a ground chart does not outvote the rest. Residuals below -12 are dropped, and so are stale bests: an official play count of at most 2 and a gap below the curve of more than $\max(6, 3\sigma)$.
2. **Fit.** All tags are fitted together by [ridge regression](#), with penalty 10 per tag, which shrinks a tag on n charts by about $n / (n + 10)$. The fit holds fixed an intercept, whether the row is a play, $\log(1 + \text{plays})$, the constant, and whether an editor tagged the chart at all. A tag needs at least 8 charts to be fitted.
3. **Chance test.** A [permutation test](#) shuffles the tags across charts 400 times. A trait's p-value is $(k + 1) / 401$, where k is the number of shuffles whose offset is at least as large.
4. **Stability test.** The player's charts are split in half at random three times, and each half is refitted. The sign must hold in both halves every time, with a coefficient of at least 0.1.

A trait is *confirmed* when $p \leq 0.02$, both halves agree, it covers at least 12 charts and its offset is at least 0.4. Chart type, era, genre and designer are measured but never reported as skills. A skill group that is not confirmed but has an offset of at least 0.25, $p \leq 0.15$ and at least 12 charts is shown as a *leaning*, and leanings are never used by the model.

The note-type losses of Section 3.12 form a second family of traits with the same shape. For each type, the offset is $-(\text{the type's loss rate} - \text{the player's overall rate}) \times \text{the type's points at stake per play}$. A type needs 200 notes. It is confirmed at 25 plays with an offset of at least 0.4, and shown as a leaning from 10 plays at 0.25.

3.15 Forecast and route planning

The rating forecast ([insights/forecast.py](#)) takes the stored rating readings from the last 120 days. It needs at least 3 readings spanning at least 10 days. Its slope is the median over every pair of readings at least a day apart, a [Theil-Sen estimator](#), which steps over the long flat stretches between sessions. The days to the next thousand are $(\text{goal} - \text{current}) / \text{slope}$, dropped if more than 730. A slope under 0.15 per day is called flat.

The route planner ([planning/build.py](#)) builds a path to a goal, by default the next thousand, greedily. It lists every (chart, target rank) pair whose target is above the current score, at most $\mu + p\sigma$ and at most 8 points above the current score, passes the S-or-better proof of Section 3.9, and has odds of at least a floor. Both p and the floor depend on the level: p is 0.4, 1.0, 1.8 or 2.6 and the floor 0.35, 0.2, 0.1 or 0.04, from Easier to Long shots. Unplayed charts from reach -1 to the search ceiling $+0.3$ enter at S or better, priced with a spread of at least $\max(0.5, \sigma c)$. At each step the planner takes the option with the largest $\text{gain} \times \text{odds} \times \text{affinity}$, counting gain against pools that already include the earlier steps, and it discounts unplayed charts by 0.7. It stops at the goal or after 40 steps. If it is still short, it may add up to 8 more never-played charts from a wider window, with a lower odds floor, and marks them.

4. Worked Example

All data in this section is an example from one Rasmai user, and it illustrates the method only. Mileage will vary from player to player: another player's curve, spreads, calibration, picks and traits will come out differently. The section runs the methods of Section 3 on that user's export, `rasmai-scores-20260930-0606.json`, recorded on 30 September 2026 at 06:06. Every figure below comes from calling the project's own functions on the export. Titles are written as the export stores them, which drops the spaces from many of them.

4.1 The example user's data

Item	Value
Scored charts	554 (350 Expert, 158 Master, 32 Advanced, 14 Re:Master)
Charts flagged as current version	122 (version 26 in otoge-db's numbering)
Charts with constant 13.0 or higher	145
Total plays	593
Stored plays	168, of which 107 have a judgement page and 108 have an earlier best on record
Recent-play list	50 plays
Charts with official play counts	91
Rating readings	26, from 3 to 30 September
Rating on the site	13,580 (silver band)

Ranks are spread across the whole ladder: 41 charts at SSS+, 41 at SSS, 43 at SS+, 41 at SS, 75 at S+ and 58 at S, with the remaining 255 below S.

Scope of the reproduction. The calls below use the export's stored constants and new-chart flags, checked against otoge-db and dxrating as fetched on 30 September 2026; enrichment changed none of the 554 constants or flags. The rerun is the full profile build: calibration, trait fit and judgement profile included.

Chart notation was read from Maichart-Converts, 6,995 charts, 437 of which match this player's 554. Three things are left out. mai ノーツ could not be reached, so there are no editor tags and the notation was not checked against its published note counts. The optional beta taste model was not run. The dashboard's numbers can differ where those matter.

4.2 Rebuilding the rating

The export stores a rating for each chart. Recomputing each one with the formula of Section 3.2 from its constant and achievement matches all 554. Two rows show the floor in action:

- BreakTheSpeakers, constant 12.9, achievement 100.1567 (SSS, $k = 21.6$): $12.9 \times 1.001567 \times 21.6 = 279.08$, so 279.
- GetU ㇏ack, constant 12.9, achievement 100.0008 (SSS): $12.9 \times 1.000008 \times 21.6 = 278.64$, so 278.

The example user's 15 new-pool charts are below. They total 4,054, the cutoff (15th place) is 260, and the next chart outside the pool is worth 257.

Chart	Tier	Constant	Achievement (%)	Rank	Rating
おべんきょうたいむ	Expert	12.4	100.5784	SSS+	279
Let'sミクササイズ!!	Expert	12.4	100.6818	SSS+	279
AiAe	Expert	12.4	100.5860	SSS+	279
BreakTheSpeakers	Expert	12.9	100.1567	SSS	279
GetU ㇏ack	Expert	12.9	100.0008	SSS	278
Myosotis	Expert	12.0	100.6776	SSS+	270
テトリス	Master	12.0	100.7807	SSS+	270
零號車輛	Expert	12.8	99.7583	SS+	269
7Wonders	Expert	12.8	99.6400	SS+	269
Worlders	Master	12.8	99.6147	SS+	269
モニタリング	Master	11.8	100.6320	SSS+	265
雲外蒼電-DreamingVoltage-	Expert	12.2	100.3777	SSS	264
トレジャーガーデン	Expert	11.7	100.7938	SSS+	263
愛スクリーム!	Re:Master	13.4	97.6384	S	261
CustomizedJustice	Expert	13.0	98.8501	S+	260

The old pool holds 35 charts totalling 9,525, from FREEDOMDiVE(tpzOvercuteRemix) at 288 down to a cutoff of 263 (LostDesire, constant 11.7 at SSS+). The next chart outside the pool is also worth 263. The pools sum to $4,054 + 9,525 = 13,579$, one point under the 13,580 the site shows. The export does not say where that point comes from.

Two features matter later. The old pool spans constants 11.7 to 13.8, and 16 of its 35 charts are at SSS+, the rank with the largest coefficient (22.4). And the old pool's cutoff equals its next value, so a chart that enters at exactly 263 gains nothing.

4.3 The fitted curve

The example user's scored charts fall into half-constant buckets. The table lists the buckets at and above 12.0 with the curve the model builds from them.

Bucket	Charts	Typical score (%)	Best (%)	Curve E(c) (%)
12.0	62	100.04	100.89	99.30
12.5	59	99.63	100.72	98.59
13.0	104	97.88	100.66	96.42
13.5	65	91.91	99.23	92.21
14.0	26	89.54	97.51	89.84
14.5	4	83.53	88.05	83.83

Buckets from 10.0 to 11.5 hold 28 to 56 charts each and sit at 99.8 to 100.3 typical, and the 3 charts in the 15.0 bucket sit at 72.3. The straight-line step gives $L(c) = 111.61 - 1.326c$ and a between-chart spread σ_c of 1.87, near its 2.0 upper clamp. At $c = 13.0$ the line alone gives 94.4, well under the 97.9 this player typically scores there. The buckets around 13.0 carry a combined kernel weight of about 21, so $\beta \approx 0.88$: the local average of about 96.7 dominates, and the blend is 96.4. That is the curve value in the table. It sits below the 13.0 bucket's own 97.9 because the kernel also reaches the 13.5 bucket at 91.9.

From the curve the profile reads three landmarks: the comfort constant is 12.2 (SS expected), the reach constant is 12.8 (S expected), and the hardest chart with an S or better is 13.8. The dense ceiling is 15.0, because even the 15.0 bucket holds 3 charts, so the dense-ceiling drop-off never applies to this player. Tier offsets come out at +1.02 for Expert, +0.16 for Advanced, -2.15 for Master and -1.90 for Re:Master, so an Expert 13 is expected about 3.2 points higher than a Master 13.

The run spread comes from 81 pairs of successive plays among the 168 stored plays. The measured value, about 0.458, is blended with the 0.5 default at weight $81 / 96$, giving 0.464. That is a quarter of the between-chart spread of 1.87, which is the gap the two-spread design exists for.

4.4 Calibration

Of the example user's 108 stored plays with an earlier best on record, 92 match a chart in the database and survive the $z < -4$ cut. Section 3.11 turns them into these corrections:

Measure	Value
Robust spread of z , after shrinkage (scale s)	0.73
Shift of the curve	+0.21
Tier shift, Expert	-0.03
Tier shift, Master	+0.03
Plays within one spread (coverage)	68%
New bests predicted	39%
New bests that happened	50%

The player's runs clustered more tightly around the prediction than the model assumed, so both spreads shrink by 27%: σ_c from 1.87 to 1.37 and σ_{run} from 0.46 to 0.34. They also landed slightly higher than the curve, so every expectation rises by 0.21. The model predicted a new best on 39% of these plays, and 50% were, so it had been slightly too cautious for this player.

4.5 Pricing one chart: GIGANTØMAKHIA

The most recent play in the example user's export is GIGANTØMAKHIA (DX Expert, constant 13.0), played on 30 September at 98.8179%, a new best over 98.7526%. The game counts 6 plays of it, and 3 of them are in the recent-play list: 97.8927, 98.7526 and 98.8179, oldest first. Its notation earns seven labels: multi-touch, trills, trills on the spot, a trill over a slide, spinning round the ring, fan slides and overlapping loop slides. The model prices a replay in six steps.

1. **Rating.** Achievement 98.8179 is S+ ($k = 20.3$): $13.0 \times 0.988179 \times 20.3 = 260.8$, so 260. This is an old-pool chart and the old cutoff is 263, so it does not count toward the rating yet.
2. **Curve.** $m = 96.42$ (curve) + 1.02 (Expert offset) + 0.21 (shift) - 0.03 (Expert tier shift) = 97.62.
3. **Spread.** $\sigma = 0.339 \times \sqrt{1 + 1/6} = 0.366$, the calibrated run spread widened for a chart played six times.
4. **Centre.** The best, 98.82, is above m , so the floor applies: $\mu = 98.82 - 0.5 \times 0.366 = 98.64$. The recent-play list shows the chart rising by 0.93 from its oldest to its latest run, which adds $\min(0.4, 0.25 \times 0.93) = 0.23$. So $\mu = 98.87$.
5. **Target.** The ceiling is $\min(100.5, 98.87 + 1.0 \times \max(0.366, 1.37)) = 100.23$, which rules out SSS+. SSS has odds $\Phi((98.87 - 100.0) / 0.366) = 0.1\%$ and SS+ has $\Phi((98.87 - 99.5) / 0.366) = 4.2\%$, both under the balanced floor of 16%. SS at 99.0 has $\Phi(-0.37) = 36\%$ and passes. Its rating is $13.0 \times 0.99 \times 20.8 = 267.7$, so 267.
6. **Gain and score.** The chart is outside the old pool, so $\Delta R = 267 - 263 = 4$. With affinity 1.82 and freshness 0.45 (three recent runs), its pick score is $4 \times 0.357 \times 1.82 \times 0.45 = 1.17$.

The pick therefore reads: one run has a 36% chance of lifting the rating by 4. Without calibration, the spread would be 0.50 and the centre 98.80, which puts the same odds at 34%. Calibration narrows the spread and raises the centre, and here that adds 2 points of probability.

The table shows every rating move on the example user's list at the balanced level, all of them Expert charts, in the order the bot shows them: by ΔR , then by odds. 44 near misses follow them.

Chart	Constant	Best (%)	Target	μ (%)	Plays	ΔR	Odds
ChurrosParlor	12.5	100.3463	SSS+	100.17	8	11	18%
雲外蒼電-DreamingVoltage-	12.2	100.3777	SSS+	100.33	6	10	32%
エータ・ベータ・イータ	12.4	100.4695	SSS+	100.30	23	10	28%
TidalWave	12.3	100.2075	SSS+	100.19	17	10	19%
BigDaddy	12.0	100.2993	SSS+	100.30	25	7	28%
LatentKingdom	12.8	99.8955	SSS	99.72	25	7	21%
CustomizedJustice	13.0	98.8501	SS	98.67	10	7	18%
Cryptarithm	13.0	99.9080	SSS	99.91	23	6	40%
GIGANTØMAKHIA	13.0	98.8179	SS	98.87	6	4	36%
AzureVixen	13.2	99.4861	SS+	99.31	13	4	30%
鬼女紅妖	11.7	100.3134	SSS+	100.31	2	3	33%
DANGEROOOOJSJUNGLE	11.8	100.4519	SSS+	100.45	2	2	45%
larva	12.9	98.8626	SS	98.69	12	2	19%

The 13 moves add 83 points if each is counted against today's pools. Landed one after another, they add 76, because each landing raises a cutoff that the next gain is measured against. The list holds no unplayed charts: no chart in the search window had a first-pass estimate of 94 or more that would enter a pool.

CustomizedJustice shows why pool position matters. It sits in the new pool at the cutoff of 260, so raising it to 267 lifts the pool by the full 7. GIGANTØMAKHIA makes the same 260 to 267 jump from outside the old pool, so it gains only the 4 above that pool's cutoff.

4.6 Where the points went

The example user's GIGANTØMAKHIA play has a judgement page. Its counts are below, with the losses from Section 3.12. The chart holds 861 weight units ($441 + 2 \times 63 + 3 \times 62 + 43 + 5 \times 13$), so one unit is $b = 100 / 861 = 0.1161$.

Note type	Notes	Great	Good	Miss	Loss (points)
Tap	441	43	7	3	1.754
Hold	63	2	0	0	0.093
Slide	62	0	0	0	0
Touch	43	0	0	0	0
Break (base)	13	1	0	0	0
Break bonus, 7 perfects and 1 great					0.315
Remainder the counts cannot price					0.020
Total					2.182

Tap notes lose $(43/5 + 7/2 + 3) \times 0.1161 = 1.754$. Counts alone explain 2.162 points, which would put the score at $101 - 2.162 = 98.838$, 0.020 above the recorded 98.8179. That 0.020 is what the page's counts cannot price, such as how much the break great cost from its base score, and the code assigns it to breaks. The total of $2.182 = 101 - 98.8179$ therefore holds by construction. It is not an independent check; the remainder shows how much the reconstruction had to assume.

On this one play, taps carry 80% of the loss on 51% of the points at stake ($441 / 861$), a tilt of 1.57. Timing shows 35 fast and 18 late, so 34% of timing errors were late. One play is not a profile, though, and the example user's 107 stored judgement pages say something different:

Note type	Notes	Share of points at stake	Share of loss	Tilt	Critical or perfect
Tap	45,696	43.7%	40.2%	0.92	86.4%
Hold	6,171	11.6%	9.3%	0.81	89.0%
Slide	5,549	16.0%	12.3%	0.77	90.3%
Touch	4,335	4.2%	4.4%	1.06	89.6%
Break	5,727	24.6%	23.1%	0.94	88.5%

The loss shares do not reach 100% because the break bonus, 10.7% of all loss, is reported apart. No type comes near the 1.5 tilt that would name a weakness, so none is named. Across all 107 plays the user lost 10.05 points a play, 1.07 of it the critical-only break bonus. Late hits make up 51.5% of timing errors (1,865 late against 1,757 fast). As traits (Section 3.14), the note types sit between -0.12 (tap) and $+0.20$ (slide) points a play, all inside the 0.25 needed for a leaning.

4.7 Traits, including simai measures

The notation of 437 of the example user's 554 charts was read, 240 of the 299 at level 12 and above. With those labels, the trait fit measures 57 groups, the 5 note types included. Six of them are leanings, the ones the Traits view shows as unconfirmed:

Group	Kind	Offset (points)	Charts	Plays	p
reaching across the screen	simai	−0.89	62	19	0.017
slides fired from one spot	simai	−0.84	45	19	0.030
trills	simai	−0.75	79	45	0.080
long streams	simai	−0.70	44	17	0.080
bursts	simai	+0.90	34	29	0.067
a slide traced straight back	simai	+0.94	63	30	0.022

None is confirmed. Reaching across the screen passes the chance test ($p = 0.017$) but not the split-half stability test, and the others miss the chance test. All six come from the notation, so without it this user would have no skill leanings at all. Four groups do pass both tests: charts by 緑風 犬三郎 (+1.52, 17 charts), BUDDiES and newer songs (+1.03, 287 charts), maimai-era songs before DX (−1.30, 61 charts) and charts by 鳩ホルダー (−1.39, 24 charts). Designer and era are not skills, so they are measured but never reported as traits.

4.8 Progress, forecast and route

The example user's stored readings run from 13,478 on 3 September to 13,580 on 30 September, a gain of 102 points over 27 days and 38 plays (555 to 593). The rating sat at 13,551 from 7 to 22 September, then rose to 13,559 on the 24th and to 13,570 and 13,580 on the 30th. The forecast of Section 3.15 returns 1.79 points per day, 12.5 per week or 54 per month. At that rate the next milestone, 14,000 (gold), lands around 22 May 2027, 235 days after the recording. The window is short, 26 readings over 26 days, so this is a rough guide.

The route planner at the balanced level agrees that 14,000 is far off. Its route takes 34 steps, 28 of them on charts the user has never played, including the 8 added at the end. It adds 226 rating at the listed targets and still falls 194 short. Weighted by each step's odds, the route is worth about 53 points. The next thousand waits on the skill curve itself rising, not on any set of targets available today.

5. Discussion

5.1 Design choices

Conservative centres. The model would rather under-promise than over-promise. The centre of a prediction never sits above a proven score except through a recent upward trend, capped at 0.4. The curve is capped by what the player has scored nearby, and S-or-better targets need a score like it on a chart at most 0.7 constants easier. Each rule exists because a more optimistic version made targets look cheap and then missed on held-out plays.

Two spreads. Separating the between-chart spread (1.37 after calibration in the example) from the run-to-run spread (0.34) is the central statistical decision. One number cannot serve both: the first is how much charts differ from each other, the second how much one run differs from the next. Pricing a replay with the first made every rank look reachable.

Robust statistics and shrinkage. Medians, quartiles and median absolute deviations resist abandoned runs. Every estimate that rests on a small sample, including tier offsets, run spread, calibration and traits, is pulled toward a neutral value by a factor of the form $n / (n + k)$. A new player sees a cautious model, and a long-time player sees one that follows their data.

Reading the notes. Editors' tags cover about one chart in ten. Measures read from simai notation reach every chart that has a file, and each named measure must rise with level and agree with the editors where they have tagged. In the example, every skill leaning came from the notation.

Tests against chance. A trait is named only if it beats shuffled tags and keeps its sign in both halves of the player's own charts. Most players see few confirmed traits at first; the example user has none. This is intended: a trait shown for every player would carry no information.

Recorded plays as a check. Because the bot keeps every play it sees, each future play tests an earlier prediction. Calibration, coverage and the predicted-against-actual new-best rate turn that into a running score for the model. For the example user it found runs tighter than assumed and new bests more common than predicted, 50% against 39%.

5.2 Limitations

- **External data.** Constants, versions, regions, notation and tags come from community sources. A wrong or missing constant changes a chart's rating and every number derived from it. The middle-of-band estimate for a missing chart is a fallback, not a measurement.
- **Normal model.** Run scores are treated as normally distributed around μ , and predictions are capped at 100.5, so the tail near the cap is only approximate.
- **One run, not many.** The odds are for a single run. A player who plays a chart ten times has better odds of reaching the target at least once, which the odds do not capture.
- **Thin data.** Charts above the player's dense ceiling rest on few scores, and the model widens its spread there rather than pretending to know more.
- **Judgement grades.** The break-bonus arithmetic assumes what a plain perfect keeps, and the page does not show finer grades, so part of each break's loss is inferred.
- **The reproduction here.** Section 4 runs the full profile build, but without maiノーツ's editor tags, without its check on notation note counts, and without the optional beta taste model. The dashboard's numbers can differ where those matter.
- **A one-point gap.** The recomputed best-50 (13,579) differs from the site's 13,580 by one. The export does not identify the cause.

Several figures come from the project's own backtests as recorded in its source comments and were not rerun for this paper: the headroom result (64% against 61%), the first-pass scatter (1.8 times wider), the loose-match coverage (95% against 62%), and the level correlations of the simai measures.

6. Conclusion

Rasmai builds on one exact fact, the game's rating formula, and estimates everything else from the player's own history. The formula fixes what a score is worth, and the two-pool total fixes which scores count. A skill curve fitted to the player's scores says what they reach at each difficulty. Two separate spreads say how sure the model is about a chart it has never seen and about one it knows, and calibration against recorded plays keeps both honest. Odds times best-50 gain then ranks every chart by what a replay is worth. Notation read from simai extends the trait analysis from the tenth of charts editors have tagged to most of the game. In the example user's data, the method reproduces all 554 chart ratings, rebuilds the best-50 to within a point, and prices a replay of GIGANTØMAKHIA at 36% for +4 rating. Its note-level accounting explains all but 0.020 of the 2.18 points lost on that play.

References

1. Cardin Nguyen (Ne-k), [Rasmai](#), source code (GPL-3.0); site at [rasmai.lol](#).
2. SEGA, [maimai DX](#) and [maimai DX NET \(international\)](#).
3. zvuc, [otoge-db](#): chart constants, versions and regions.

4. gekichumai, [dxrating](#): the `dxdata.json` file of revised constants, per-chart regions and aliases.
5. arussin, [maimai-chart-browser](#): search aliases and two slide-pattern rules.
6. Neskol, [Maichart-Converts](#): simai notation used in Section 4.
7. [maiノーツ](#): editors' pattern tags and chart notation.
8. SilentBlue RemyWiki, "maimai DX: Rating": rating colour bands.
9. SpiritsUnite, [maimai-score-details](#) (Apache-2.0): note-loss arithmetic.
10. Wikipedia, background on [weighted least squares](#), [kernel regression](#), [median absolute deviation](#), [ridge regression](#), [permutation tests](#) and the [Theil-Sen estimator](#).

The licences and exact file paths of the third-party sources are listed in the repository's [THIRD_PARTY_NOTICES.md](#).

Appendix A. Source map

Topic	Section	File
Rank table, coefficients, rating formula, challenge levels, colour bands	3.1, 3.2, 3.5, 3.9	engine/analysis/rating.py
Chart index, constants, level bands	3.3	engine/analysis/charts.py
Enrichment, best-50 pools, marginal gain	3.3, 3.4	engine/analysis/pools.py
Skill-curve fit	3.6	engine/analysis/profile/fit.py
Curve, spreads, chart expectation, dropped bests	3.6 to 3.8	engine/analysis/profile/model.py
Profile assembly, tier offsets, current version	3.3, 3.6	engine/analysis/profile/build.py
Target ranks and freshness	3.9, 3.10	engine/analysis/picks/targets.py
Pick ordering and near misses	3.10	engine/analysis/picks/generate.py
Unplayed charts	3.10	engine/analysis/picks/candidates.py
Run spread and calibration	3.7, 3.11	engine/insights/calibration.py
Note-level losses	3.12	engine/losses.py
Judgement profile and note-type traits	3.12, 3.14	engine/judgements.py
simai parsing	3.13	engine/simai/parse.py
simai measures, labels and thresholds	3.13	engine/simai/features.py
simai technique detectors	3.13	engine/simai/techniques.py
Notation reading, caching and matching	3.13	scraping/simai.py , scraping/simai_bulk.py
Chart attributes used as trait tags	3.14	engine/insights/tags.py
Chart traits	3.14	engine/insights/traits.py
Rating forecast	3.15	engine/insights/forecast.py
First-pass prediction and rank ladder	3.8	engine/analysis/planning/predict.py
Route planner	3.15	engine/analysis/planning/build.py
Replaying an export through the model	4	tools/replay.py